



# Recovery: ARIES normal operation and crash recovery procedure

## Wrap-up of first part of course

ACS, Marcos Vaz Salles

### **NOTE: Transactional Scale-Up Experiments**

- Refer to Part 12 of compendium, DeWitt and Gray paper: transactional scale-up consists of **"N-times as many clients, submitting N-times as many requests against N-times larger database."**
- How to operationalize evaluation?
  - Find configuration for unit system size (e.g., one core)
  - Scale system size together with configuration

# Do-it-yourself-recap: What's Stored Where

Explain each of the ARIES structures below to your colleague!



## LogRecords

prevLSN  
XID  
type  
pageID  
length  
offset  
before-image  
after-image



## Data pages

each  
with a  
pageLSN

**master record**



## Xact Table

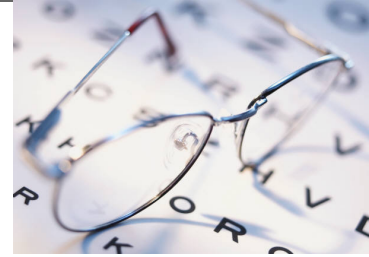
lastLSN  
status

## Dirty Page Table

recLSN

**flushedLSN**

# What should we learn today?



- Explain how write-ahead logging is achieved in the ARIES protocol
- Explain the functions of recovery metadata such as the transaction table and the dirty page table
- Interpret the contents of the log resulting from ARIES normal operation
- Explain the three phases of ARIES crash recovery: analysis, redo, and undo
- Predict how recovery metadata, system state, and the log are updated during recovery

+ a discussion of where we got so far, if time allows

## Normal Execution of an Xact

Keep in Mind:  
It must be OK to  
crash at **any time**  
→ **repeat history!**

- Series of **reads** & **writes**, followed by **commit** or **abort**.
  - We will assume that write is atomic on disk.
    - In practice, additional details to deal with non-atomic writes.
- **Strict 2PL** → concurrency is correctly handled
- **STEAL, NO-FORCE** buffer management, with **Write-Ahead Logging**.



# Transaction Commit

- Write **commit** record to log.
- All log records up to Xact's **lastLSN** are flushed.
  - Guarantees that **flushedLSN**  $\geq$  **lastLSN**.
  - Note that log flushes are sequential, synchronous writes to disk.
  - Many log records per log page.
- Commit() returns.
- Write **end** record to log.

Why?

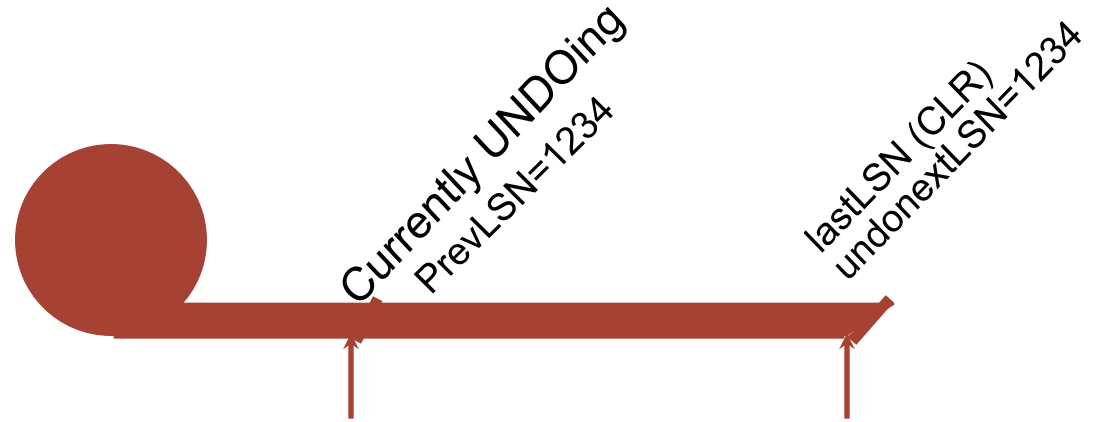


# Simple Transaction Abort

- For now, consider an explicit abort of a Xact.
  - No crash involved.
- We want to “play back” the log in reverse order, UNDOing updates.
  - Get **lastLSN** of Xact from Xact table.
  - Can follow chain of log records backward via the **prevLSN** field.
  - Before starting UNDO, write an **Abort log record**.
    - For recovering from crash during UNDO!



## Abort, cont.



- To perform UNDO, must have a lock on data!
  - Strict 2PL enforces this
- Before restoring old value of a page, write a CLR:
  - You continue logging while you UNDO!!
  - CLR has one extra field: **undonextLSN**
    - Points to the next LSN to undo (i.e. the prevLSN of the record we're currently undoing).
  - CLR *never* Undone (but they might be Redone when repeating history: guarantees Atomicity!)
- At end of UNDO, write an **end** log record.



# Checkpointing

- Periodically, the DBMS creates a checkpoint, in order to minimize the time taken to recover in the event of a system crash. Write to log:
  - **begin\_checkpoint** record: Indicates when chkpt began.
  - **end\_checkpoint** record: Contains current *Xact table* and *dirty page table*. This is a **'fuzzy checkpoint'**:
    - Other Xacts continue to run; so these tables accurate only as of the time of the **begin\_checkpoint** record.
    - **No attempt to force dirty pages to disk**; effectiveness of checkpoint limited by oldest unwritten change to a dirty page, **minDirtyPagesLSN**.
    - Use **background process** to flush dirty pages to disk!
  - Store LSN of chkpt record in a safe place (**master** record).





## Example

- 10 T1 writes P5
- 20 T2 writes P17
- 30 T1 writes P3

P3 written to disk

(pageLSN for page 3 at this time is 30)

- 40 T1 aborts
- 50 CLR T1 P3 (undonextLSN: 10)
- 60 CLR T1 P5 (undonextLSN: NULL)
- 70 End T1



## A Longer Example

- 10 T1 writes P3 (prevLSN: NULL)
- 20 T2 writes P4 (prevLSN: NULL)
- 30 T2 writes P5 (prevLSN: 20)
- flushedLSN = 20

P4 gets written to disk (pageLSN for page 4 = 20)

T2 aborts

- 50 Abort T2
- 60 CLR T2 P5 (undoNextLSN = 20), pageLSN(P5)=60

Update P5 in the buffer manager

Flush log up to log record 60

Buffer manager writes P5 to disk.

- 70 CLR T2 P4 (undoNextLSN = NULL), pageLSN(P4)=70

Update page P4

- 80 End T2
- 90 T1 commits

Flush log up to log record 90, then the commit(T1) returns

Discussion: Does this example make sense?  
Can you explain to your colleague what happened?

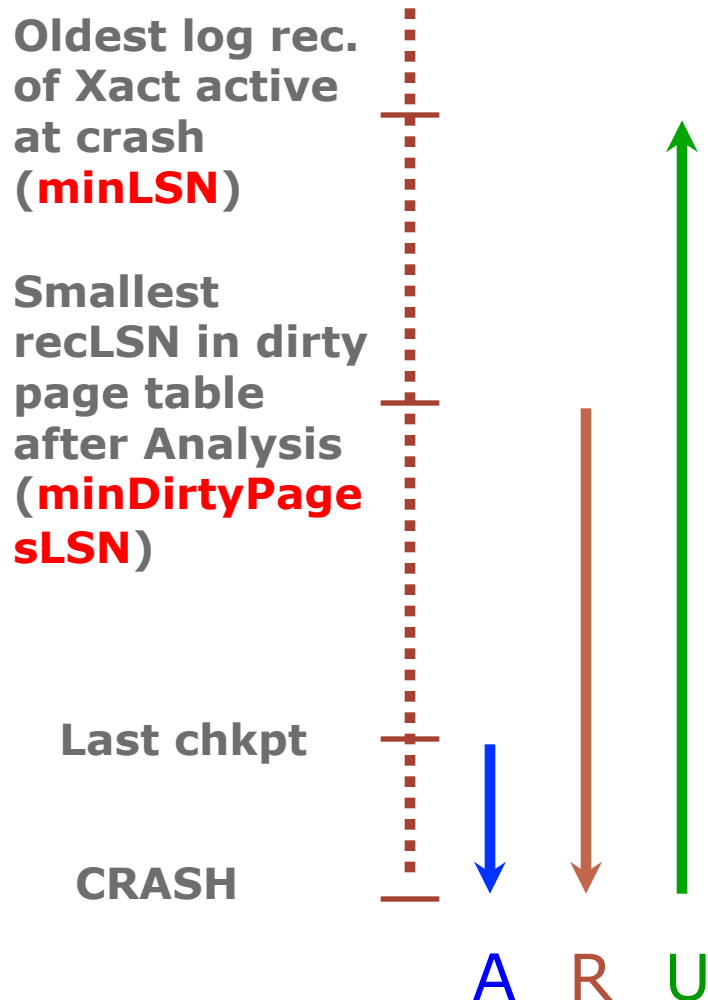


Questions so far?



# Crash Recovery: Big Picture

Keep in Mind:  
It must be OK to  
crash at **any time**  
(including during  
recovery)



- ❖ Start from a **checkpoint** (found via **master** record).
- ❖ Three phases. Need to:
  - Figure out which Xacts committed since checkpoint, which failed (**Analysis**).
  - **REDO** *all* actions.
    - ◆ **Repeat History**
  - **UNDO** effects of failed Xacts.

## Recovery: The Analysis Phase

- Reconstruct state at checkpoint.
  - via **end\_checkpoint** record.
- Scan log forward from checkpoint.
  - **End** record: Remove Xact from Xact table.
  - **Other records**: Add Xact to Xact table, set **lastLSN=LSN**, change Xact status on **commit**.
  - **Update** record: If P not in Dirty Page Table,
    - Add P to D.P.T., set its **recLSN=LSN**.



## Recovery: The REDO Phase

- We *repeat History* to reconstruct state at crash:
  - Reapply *all* updates (even of aborted Xacts!), redo CLR's.
- Scan forward from log rec containing smallest **recLSN** in D.P.T. For each CLR or update log rec **LSN**, REDO the action unless:
  - Affected page is not in the Dirty Page Table, or
  - Affected page is in D.P.T., but has **recLSN** > **LSN**, or
  - **pageLSN** (in DB) >= **LSN**.
- To **REDO** an action:
  - Reapply logged action.
  - Set **pageLSN** to **LSN**. No additional logging! (**Why?**)



## Recovery: The UNDO Phase

$ToUndo = \{ lsn \mid lsn \text{ a lastLSN of a "loser" Xact} \}$

### Repeat:

- Choose largest LSN among ToUndo.
- If this LSN is a CLR and  $undonextLSN == NULL$ 
  - Write an End record for this Xact.
- If this LSN is a CLR, and  $undonextLSN \neq NULL$ 
  - Add  $undonextLSN$  to ToUndo
- Else this LSN is an update. Undo the update, write a CLR, add  $prevLSN$  to ToUndo.

**Until ToUndo is empty.**



# Example of Recovery



Xact Table

lastLSN

status

Dirty Page Table

recLSN

flushedLSN

ToUndo

LSN	LOG
00	begin_checkpoint
05	end_checkpoint
10	update: T1 writes P5
20	update T2 writes P3
30	T1 abort
40	CLR: Undo T1 LSN 10
45	T1 End
50	update: T3 writes P1
60	update: T2 writes P5
X CRASH, RESTART	

prevLSNs



# Example: Crash During Restart!

Another  
example: crash  
during recovery!



Xact Table

lastLSN

status

Dirty Page Table

recLSN

flushedLSN

ToUndo

LSN	LOG
00,05	+ begin_checkpoint, end_checkpoint
10	+ update: T1 writes P5
20	+ update T2 writes P3
30	+ T1 abort
40,45	+ CLR: Undo T1 LSN 10, T1 End
50	+ update: T3 writes P1
60	+ update: T2 writes P5
	✗ CRASH, RESTART
64,65,70	+ T2 abort, T3 abort, CLR: Undo T2 LSN 60
80,85	+ CLR: Undo T3 LSN 50, T3 end
	✗ CRASH, RESTART
90, 95	+ CLR: Undo T2 LSN 20, T2 end

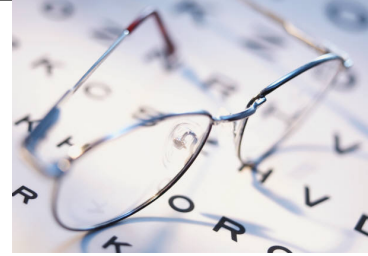
undonextLSN

## Additional Crash Issues

- What happens if system crashes during Analysis?  
During REDO?
- How do you limit the amount of work in REDO?
  - Flush asynchronously in the background.
  - Watch “hot spots”!
- How do you limit the amount of work in UNDO?
  - Avoid long-running Xacts.



# What should we learn today?



- Explain how write-ahead logging is achieved in the ARIES protocol
- Explain the functions of recovery metadata such as the transaction table and the dirty page table
- Interpret the contents of the log resulting from ARIES normal operation
- Explain the three phases of ARIES crash recovery: analysis, redo, and undo
- Predict how recovery metadata, system state, and the log are updated during recovery

+ a discussion of where we got so far, if time allows

## ACS: Evaluating the Course

- Your qualitative feedback is very important!



Comments on the course?  
Syllabus?  
Lectures?  
TA sessions? Assignments?

## Homework discussion: Granularity of Locking

- Suppose application with operations that **read on a predicate on whole database**
  - E.g., Programming Assignment 2! 😊
- **Locking Approach 1:**  
Single database lock
- **Locking Approach 2:**  
Multi-granularity locking
  - For example, one database lock + individual locks per item
- What are the trade-offs between the two approaches?
- Compare them in terms of:
  - (a) Overhead
  - (b) Implementation complexity
  - (c) Concurrency

	--	IS	IX	S	X
--	✓	✓	✓	✓	✓
IS	✓	✓	✓	✓	
IX	✓	✓	✓		
S	✓	✓		✓	
X	✓				