

ADVANCED JAVA

INTRODUCTION TO CONCURRENCY

Vivek Shah bonii@di.ku.dk

August 20, 2018

DIKU, University of Copenhagen

Concurrency Primer

Thread Creation

Thread Waiting and Interruption

High-Level ThreadPool Abstractions

CONCURRENCY PRIMER

WHAT IS CONCURRENCY?

WHAT IS CONCURRENCY?

- Appear to run several programs or several parts of a program simultaneously.

WHAT IS CONCURRENCY?

- Appear to run several programs or several parts of a program simultaneously.
- Historical progression of concurrency:

WHAT IS CONCURRENCY?

- Appear to run several programs or several parts of a program simultaneously.
- Historical progression of concurrency:
 1. Batch Processing (Hardware). Separate programs.

WHAT IS CONCURRENCY?

- Appear to run several programs or several parts of a program simultaneously.
- Historical progression of concurrency:
 1. Batch Processing (Hardware). Separate programs.
 2. Processes (OS). Separate address spaces.

WHAT IS CONCURRENCY?

- Appear to run several programs or several parts of a program simultaneously.
- Historical progression of concurrency:
 1. Batch Processing (Hardware). Separate programs.
 2. Processes (OS). Separate address spaces.
 3. Threads (Application). Same address space.

WHAT IS CONCURRENCY?

- Java is a multi-threaded language
 - Thread support for sequential programs.
 - Well-defined memory model.

WHAT IS CONCURRENCY?

- Java is a multi-threaded language
 - Thread support for sequential programs.
 - Well-defined memory model.
- Advantages of using Threads:
 1. Light-weight.
 2. Flexibility.
 3. Ease of resource sharing.

WHAT IS CONCURRENCY?

- Java is a multi-threaded language
 - Thread support for sequential programs.
 - Well-defined memory model.
- Advantages of using Threads:
 1. Light-weight.
 2. Flexibility.
 3. Ease of resource sharing.
- Disadvantages of using Threads:
 - With great powers come great responsibilities.

1. Multi-processor system

- Faster execution.

1. Multi-processor system

- Faster execution.

2. Single processor system

- Faster execution/performance.
- Sounds counter intuitive. What about the overheads of context switches ?

1. Multi-processor system

- Faster execution.

2. Single processor system

- Faster execution/performance.
- Sounds counter intuitive. What about the overheads of context switches ?
- Waiting
 - For Memory, Disk, Network or something else.
 - For responsiveness in user interactive interfaces (Event driven programming).

1. Multi-processor system

- Faster execution.

2. Single processor system

- Faster execution/performance.
- Sounds counter intuitive. What about the overheads of context switches ?
- **Waiting**
 - For Memory, Disk, Network or something else.
 - For responsiveness in user interactive interfaces (Event driven programming).
- **Improving code design**
 - Modeling simulations (games).
 - No need to worry about scheduling.
 - Modeling distributed systems (e.g., actor model)

THREAD CREATION

The `java.lang.Runnable` interface:

```
public interface Runnable {  
    void run();  
}
```

- A simple interface to a runnable piece of code.
- A thread's execution entry point is specified using `Runnable`.

```
class Thread implements Runnable {  
    Thread(Runnable target)  
    static Thread currentThread()  
    void run()  
    void start()  
    static void sleep(long millis) throws InterruptedException  
    static void yield()  
    void join() throws InterruptedException  
    void interrupt()  
    boolean isInterrupted()  
    static boolean interrupted()  
}
```

Method 1: Implement the `java.lang.Runnable` interface.

```
public class HelloWorldRunnable implements Runnable {  
    public void run() {  
        System.out.println("Hello World");  
    }  
  
    public static void main(String[] args) {  
        new Thread(new HelloWorldRunnable()).start();  
    }  
}
```

Method 2: Extend the `java.lang.Thread` class.

```
public class HelloWorldThread extends Thread {  
    public void run() {  
        System.out.println("Hello World");  
    }  
  
    public static void main(String[] args) {  
        new HelloWorldThread().start();  
    }  
}
```

- A thread performs the task defined in `run` method.
- implement `java.lang.Runnable` interface or extend `java.lang.Thread` class

- A thread performs the task defined in `run` method.
- implement `java.lang.Runnable` interface or extend `java.lang.Thread` class
 - Implementing `java.lang.Runnable` provides purer composition.
 - Extending `java.lang.Thread` is useful for simple applications but limited owing to sub-classing.

THREAD WAITING AND INTERRUPTION

- `Thread.sleep()` and `Thread.join()` cause the **current thread** to wait.
- `Thread.sleep(x)` - Wait for **x** milliseconds.
 - Why is sleep static?

- `Thread.sleep()` and `Thread.join()` cause the **current thread** to wait.
- `Thread.sleep(x)` - Wait for **x** milliseconds.
 - Why is sleep static?
 - Because a thread cannot put another thread to sleep.
- `t.join()` - Wait until **thread t** dies.

- Calling `x.interrupt()` signals **thread x** to stop doing what it is doing and do something else. (e.g., terminate).
- From the JavaDoc:
"If this thread is blocked in an invocation of ... `join()` ... or `sleep(long)`, then its interrupt status will be cleared and it will receive an `InterruptedException`."
...
"If none of the previous conditions hold then this thread's interrupt status will be set."

- What happens when a thread is interrupted ?
 - If the thread is waiting, `java.lang.InterruptedException` is thrown.
 - Otherwise, nothing happens.
 - Unless the thread inspects its interrupted status regularly.

- What happens when a thread is interrupted ?
 - If the thread is waiting, `java.lang.InterruptedException` is thrown.
 - Otherwise, nothing happens.
 - Unless the thread inspects its interrupted status regularly.
- Good practice:
 - Do not use `while(true)` to make a thread run forever.
 - Use `while(!Thread.interrupted())` instead.

- `Thread.yield()` signals the scheduler to schedule another thread for execution. *Just a hint.*

- `Thread.yield()` signals the scheduler to schedule another thread for execution. *Just a hint.*
- `Thread.setPriority()` can set the scheduling priority of a thread.
 - `MAX_PRIORITY`, `MIN_PRIORITY`, `NORM_PRIORITY`.

- `Thread.yield()` signals the scheduler to schedule another thread for execution. **Just a hint.**
- `Thread.setPriority()` can set the scheduling priority of a thread.
 - `MAX_PRIORITY`, `MIN_PRIORITY`, `NORM_PRIORITY`.
- User Threads and Daemon Threads.
 - JVM runs a process until at least one user thread is running.
 - `Thread.setDaemon(boolean)` function can turn on/off the daemon status of a thread.

- Sending and handling Interrupts.

- Sending and handling Interrupts.
- `t.interrupt()` sends an interrupt to thread `t` for handling.
Thread `t` can detect the interrupt by

- Sending and handling Interrupts.
- `t.interrupt()` sends an interrupt to thread `t` for handling. Thread `t` can detect the interrupt by
 - Catching `java.lang.InterruptedException`.
 - If you do nothing, set the status back.
 - `Thread.currentThread().interrupt()`.

- Sending and handling Interrupts.
- `t.interrupt()` sends an interrupt to thread `t` for handling. Thread `t` can detect the interrupt by
 - Catching `java.lang.InterruptedException`.
 - If you do nothing, set the status back.
 - `Thread.currentThread().interrupt()`.
 - Checking status using `Thread.interrupted()`. Do not check after catching an `InterruptedException`.

- Sending and handling Interrupts.
- `t.interrupt()` sends an interrupt to thread `t` for handling. Thread `t` can detect the interrupt by
 - Catching `java.lang.InterruptedException`.
 - If you do nothing, set the status back.
 - `Thread.currentThread().interrupt()`.
 - Checking status using `Thread.interrupted()`. Do not check after catching an `InterruptedException`.
- Interrupts are used to cancel a thread. Good programs support thread cancellation.

- Sending and handling Interrupts.
- `t.interrupt()` sends an interrupt to thread `t` for handling. Thread `t` can detect the interrupt by
 - Catching `java.lang.InterruptedException`.
 - If you do nothing, set the status back.
 - `Thread.currentThread().interrupt()`.
 - Checking status using `Thread.interrupted()`. Do not check after catching an `InterruptedException`.
- Interrupts are used to cancel a thread. Good programs support thread cancellation.
- Let us see thread waiting and manipulation in action.

WHAT CAN WE DO (SAFELY)?

- Concurrency with no shared memory.
 - No communication (besides interruptions).
 - Data parallelism.

WHAT CAN WE DO (SAFELY)?

- Concurrency with no shared memory.
 - No communication (besides interruptions).
 - Data parallelism.
- Parallel version of quicksort.
- Parallel linear search.
- Parallel binary search.

WHAT CAN WE DO (SAFELY)?

- Concurrency with no shared memory.
 - No communication (besides interruptions).
 - Data parallelism.
- Parallel version of quicksort.
- Parallel linear search.
- Parallel binary search.
- Algorithms where threads can work independently.

HIGH-LEVEL THREADPOOL ABSTRACTIONS

- Thread creation and destruction are expensive.
- Large scale applications require separate thread management and creation.
- `java.util.concurrent` package provides thread management via **Executors**.
- Executors are a layer of indirection between client and execution of a task.
- Multiple types of threadpools like `CachedThreadPool`, `FixedThreadPool`.
- Frees the application from thread-management and focus on resource sharing and co-ordination among threads.

- `java.util.concurrent.Callable` is an extension of the `java.lang.Runnable` interface which represents a task that returns a result or throws an exception.
- Must implement `call()` method instead of `run()` which returns a result.

- `java.util.concurrent.Callable` is an extension of the `java.lang.Runnable` interface which represents a task that returns a result or throws an exception.
- Must implement `call()` method instead of `run()` which returns a result.
- Is only supported for execution through `java.util.concurrent.Executor` interface.
- Submitting a task obtains a `java.util.concurrent.Future` object.

- Calling `get()` on the object blocks if the thread has not completed.
 - Timed version of `get()` is available.
 - `isDone()` is a non blocking version which can check to see if the result is available.
- Calling `cancel(boolean)` interrupts the thread executing the task.

- Calling `get()` on the object blocks if the thread has not completed.
 - Timed version of `get()` is available.
 - `isDone()` is a non blocking version which can check to see if the result is available.
- Calling `cancel(bool)` interrupts the thread executing the task.
- **Let us see Executors and Callable in action.**

- Threads do not always operate on independent resources or wait for other threads to finish.

- Threads do not always operate on independent resources or wait for other threads to finish.
- With concurrency resource contention and sharing is a problem that needs to be tackled.
- Proper access to resources must be ensured.

- Threads do not always operate on independent resources or wait for other threads to finish.
- With concurrency resource contention and sharing is a problem that needs to be tackled.
- Proper access to resources must be ensured.
- Mechanisms that can be used.
 - Monitors.
 - Using explicit Lock objects.
 - Atomic classes and volatility.
 - High-level synchronizers.

- Assignment 1 will be released.
- Lab sessions (work on assignment) will take place from 13:30 – 17:00.
- Work in groups of 2-3. **Inform us of your group by 14:00.**
- No need for a report, comment your code :-).
- **Deadline 23:59 hrs today.**
- **Do not hesitate, ask questions if need be.**
- Any solution that conforms to the handed out interfaces is acceptable.