

# ADVANCED JAVA

## JAVA VIRTUAL MACHINE

---

Vivek Shah [bonii@di.ku.dk](mailto:bonii@di.ku.dk)

August 20, 2018

DIKU, University of Copenhagen

Java Technology Overview


JVM Tuning Options

Profiling and Other Utilities

# JAVA TECHNOLOGY OVERVIEW


---

# HOW JAVA PROGRAMS ARE EXECUTED?



Overview of Java Development Process (Oracle Docs)


# HOW JAVA PROGRAMS ARE EXECUTED?



Overview of Java Development Process (Oracle Docs)

- Source programs are written in plain-text in `.java` files


# HOW JAVA PROGRAMS ARE EXECUTED?



Overview of Java Development Process (Oracle Docs)

- Source programs are written in plain-text in `.java` files
- Java compiler compiles them into `.class` files (consisting of bytecodes)


# HOW JAVA PROGRAMS ARE EXECUTED?



Overview of Java Development Process (Oracle Docs)

- Source programs are written in plain-text in `.java` files
- Java compiler compiles them into `.class` files (consisting of bytecodes)
- The Java Virtual Machine runs the `.class` files

# OVERVIEW OF JVM



Overview of JVM (Oracle Docs)



- Bytecodes
  - Instruction set of the JVM
  - Both a stack and register machine

- Bytecodes
  - Instruction set of the JVM
  - Both a stack and register machine
- Java Virtual Machine
  - Abstract machine that can execute Java byte-code → Java software processor
  - JVM components
    1. Bytecode verifier
    2. Memory manager
    3. Execution engine (Interpreter/JIT compiler)

Why do we need a JVM to interpret the .class file instead of compiling to machine code?

## WHY USE A JVM?

1. Platform independence → Source and compiled code is independent of hardware and operating system platform
  - Compiled code is just dependent on the JVM platform
2. Allows distribution of binaries (`.class` files)
3. Compilation process is platform independent but compilers are not always
4. Compile once, run anywhere
5. Allows sand-boxing e.g., Applets
6. Allows instrumentation. Dynamic memory management.

## WHY USE A JVM? (CONTD.)

7. Provides Extensibility
  - Any language that can compile down to bytecode can be interpreted (e.g., Scala, Kotlin)
8. Only JVM needs to be platform dependent. JVM writers can focus on efficient processing of bytecodes instead of the whole upper language stack
9. **Has overheads.** Various optimizations like JIT compilation, hardware processors supporting bytecodes

## JVM TUNING OPTIONS

---

- JVM provides a **tunable** execution environment
- Various ways to affect performance of programs by tuning JVM
  1. **Memory management** (heap tuning and garbage collection)
  2. **Code generation** (e.g., control inlining of code)
  3. **Configuration of environment variables using properties**
  4. **Usage of JVM hooks**
- Use **limited** profiling support to understand program performance
- **More than meets the eye** (Use extreme care)

### Meaning of the option prefix

- -x → non standard across different VM, subject to change without notice
- -xx → Non standard and not stable

For standard options see the reference pages for the platform



## FEW COMMONLY USED GENERIC JVM TUNING OPTIONS

Option	Meaning

## FEW COMMONLY USED GENERIC JVM TUNING OPTIONS

Option	Meaning
<code>-classpath</code>	Directories to search for class files

## FEW COMMONLY USED GENERIC JVM TUNING OPTIONS

Option	Meaning
-classpath	Directories to search for class files
-d32 or -d64	Run in 32 bit or 64 bit enviroment

## FEW COMMONLY USED GENERIC JVM TUNING OPTIONS

Option	Meaning
-classpath	Directories to search for class files
-d32 or -d64	Run in 32 bit or 64 bit enviroment
-server	Turns on optimizing JIT

## FEW COMMONLY USED GENERIC JVM TUNING OPTIONS

Option	Meaning
-classpath	Directories to search for class files
-d32 or -d64	Run in 32 bit or 64 bit environment
-server	Turns on optimizing JIT
-verbose	Runs in verbose mode

## FEW COMMONLY USED GENERIC JVM TUNING OPTIONS

Option	Meaning
-classpath	Directories to search for class files
-d32 or -d64	Run in 32 bit or 64 bit environment
-server	Turns on optimizing JIT
-verbose	Runs in verbose mode
-jar	Execute a program in a jar file

## FEW COMMONLY USED GENERIC JVM TUNING OPTIONS

Option	Meaning
-classpath	Directories to search for class files
-d32 or -d64	Run in 32 bit or 64 bit environment
-server	Turns on optimizing JIT
-verbose	Runs in verbose mode
-jar	Execute a program in a jar file
-Xint	Turns off compilation to native code, only interpretation

## FEW JVM MEMORY TUNING OPTIONS (HEAP AND STACK)

Option	Meaning



## FEW JVM MEMORY TUNING OPTIONS (HEAP AND STACK)

Option	Meaning
-Xms	Initial size of heap e.g., -Xms1024m

## FEW JVM MEMORY TUNING OPTIONS (HEAP AND STACK)

Option	Meaning
-Xms	Initial size of heap e.g., -Xms1024m
-Xmx	Maximum size of heap e.g., -Xmx8192m

## FEW JVM MEMORY TUNING OPTIONS (HEAP AND STACK)

Option	Meaning
-Xms	Initial size of heap e.g., -Xms1024m
-Xmx	Maximum size of heap e.g., -Xmx8192m
-Xss	Sets the stack size, -Xss4096k ← this value will be multiplied by the number of threads

- Tuning garbage collection is not only about setting parameters but understanding the process
- Generational garbage collection
- Can choose the type of garbage collector
- Blocking versus Incremental versus Concurrent

## FEW GARBAGE COLLECTOR TUNING OPTIONS

Option	Meaning

## FEW GARBAGE COLLECTOR TUNING OPTIONS

Option	Meaning
<code>-XX:+UseSerialGC</code>	Enables the use of the serial garbage collector (Good when running many JVM's)

## FEW GARBAGE COLLECTOR TUNING OPTIONS

Option	Meaning
<code>-XX:+UseSerialGC</code>	Enables the use of the serial garbage collector (Good when running many JVM's)
<code>-XX:+UseParallelGC</code>	Enables the use of the parallel scavenge garbage collector (also known as the throughput collector)

## FEW GARBAGE COLLECTOR TUNING OPTIONS

Option	Meaning
<code>-XX:+UseSerialGC</code>	Enables the use of the serial garbage collector (Good when running many JVM's)
<code>-XX:+UseParallelGC</code>	Enables the use of the parallel scavenge garbage collector (also known as the throughput collector)
<code>-XX:+UseParallelOldGC</code>	Enables the use of the parallel garbage collector for full GCs (heavy duty, long pauses)



## FEW GARBAGE COLLECTOR TUNING OPTIONS (CONTD.)

Option	Meaning

## FEW GARBAGE COLLECTOR TUNING OPTIONS (CONTD.)

Option	Meaning
-XX:+UseConcMarkSweepGC	Enables the use of the CMS garbage collector for the old generation (short pauses)

## FEW GARBAGE COLLECTOR TUNING OPTIONS (CONTD.)

Option	Meaning
-XX:+UseConcMarkSweepGC	Enables the use of the CMS garbage collector for the old generation (short pauses)
-XX:+UseG1GC	Enables the use of the garbage-first (G1) garbage collector (better CMS)

## FEW GARBAGE COLLECTOR TUNING OPTIONS (CONTD.)

Option	Meaning
-XX:+UseConcMarkSweepGC	Enables the use of the CMS garbage collector for the old generation (short pauses)
-XX:+UseG1GC	Enables the use of the garbage-first (G1) garbage collector (better CMS)


Use these options only if you understand the process and its impact

# HEAP STRUCTURES FOR GARBAGE COLLECTION

## Hotspot Heap Structure



## G1 Heap Allocation



- E** Eden Space
- S** Survivor Space
- O** Old Generation

## FEW GARBAGE COLLECTOR HEAP TUNING OPTIONS

Option	Meaning

## FEW GARBAGE COLLECTOR HEAP TUNING OPTIONS

Option	Meaning
-XX:NewSize	Size of new generation

## FEW GARBAGE COLLECTOR HEAP TUNING OPTIONS

Option	Meaning
-XX:NewSize	Size of new generation
-XX:MaxNewSize	Maximum size of new generation



## FEW GARBAGE COLLECTOR HEAP TUNING OPTIONS

Option	Meaning
-XX:NewSize	Size of new generation
-XX:MaxNewSize	Maximum size of new generation
-XX:NewRatio	Ratio of old/new generations

## FEW GARBAGE COLLECTOR HEAP TUNING OPTIONS

Option	Meaning
-XX:NewSize	Size of new generation
-XX:MaxNewSize	Maximum size of new generation
-XX:NewRatio	Ratio of old/new generations
-XX:SurvivorRatio	Ratio of eden/survivor space

## FEW GARBAGE COLLECTOR HEAP TUNING OPTIONS

Option	Meaning
-XX:NewSize	Size of new generation
-XX:MaxNewSize	Maximum size of new generation
-XX:NewRatio	Ratio of old/new generations
-XX:SurvivorRatio	Ratio of eden/survivor space

Lots more !! Use these options only if you understand them

## FEW CODE GENERATION TUNING OPTIONS (INLINING)

Option	Meaning

## FEW CODE GENERATION TUNING OPTIONS (INLINING)

Option	Meaning
<code>-XX:MaxInlineSize</code>	Sets the maximum size in bytes of bytecode that will be inlined

## FEW CODE GENERATION TUNING OPTIONS (INLINING)

Option	Meaning
<code>-XX:MaxInlineSize</code>	Sets the maximum size in bytes of bytecode that will be inlined
<code>-XX:FreqInlineSize</code>	Sets the maximum size in bytes of bytecode of frequently executed method that will be inlined

## FEW CODE GENERATION TUNING OPTIONS (INLINING)

Option	Meaning
<code>-XX:MaxInlineSize</code>	Sets the maximum size in bytes of bytecode that will be inlined
<code>-XX:FreqInlineSize</code>	Sets the maximum size in bytes of bytecode of frequently executed method that will be inlined
<code>-XX:LoopUnrollLimit=n</code>	Unrolls loops with a compiler internal representation node count of less than n

- Properties passed using `-Dproperty=p`
- Retrieve the property value using `System.getProperty(p)`
- Some property keys are already predefined from which certain environment values can be retrieved
- Use `.properties` file instead of long list of properties in command line



## PROFILING AND OTHER UTILITIES

---

- Profiling is done in order to understand performance trade-offs
  - Use a profiler to understand your code better
- `-Xprof` → option for some basic profiling

- Profiling is done in order to understand performance trade-offs
  - Use a profiler to understand your code better
- `-Xprof` → option for some basic profiling
- JVM Monitor has nice integration with Eclipse
- Most full featured profilers are commercial

1. jdb - Java debugger
2. javap - Java disassembler
3. javadoc - Java document generator